

## Data Assessment – 1

Tasks:

### Day 1: Setup and Data Collection

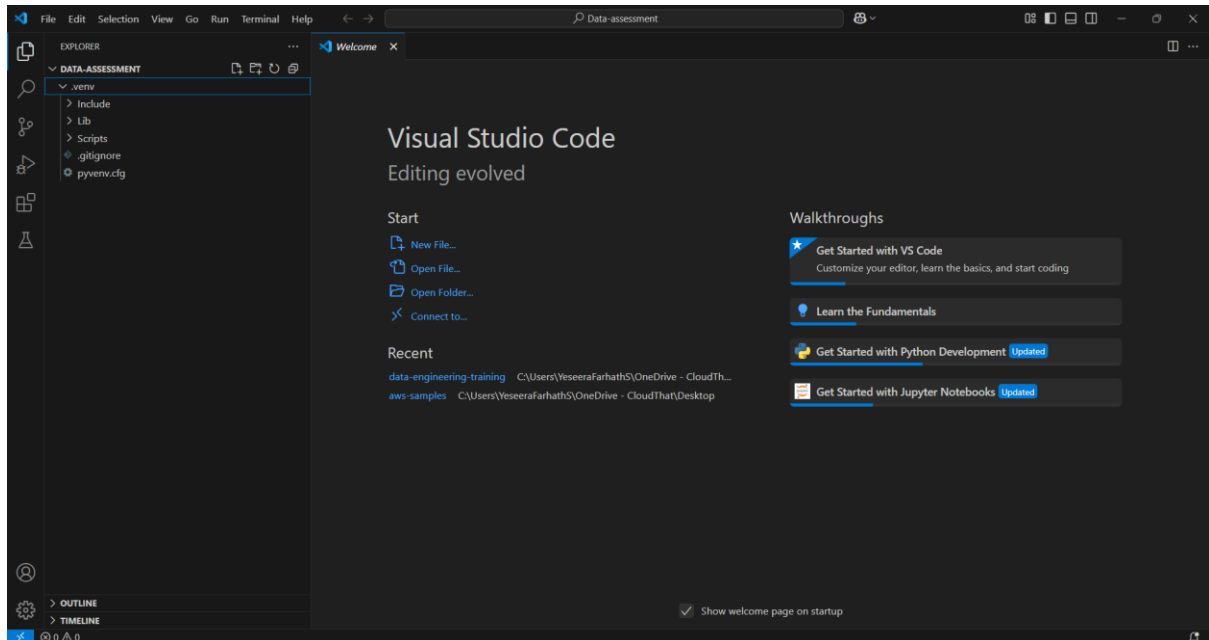
#### 1. Team Setup

- Set up virtual environments for Python dependencies
- Create a new GitHub repository for the project
- Install required Python libraries (NumPy, Pandas, Matplotlib, scikit-learn)

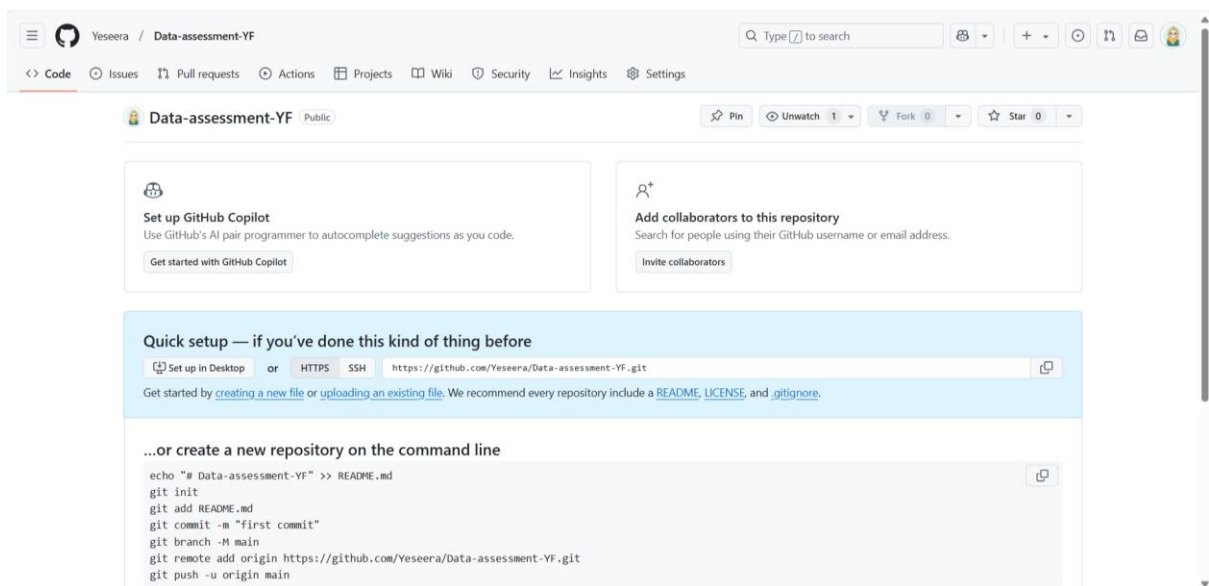
#### 2. **Dataset Selection:** Choose any from one dataset from [data.gov.in](https://data.gov.in)

#### 3. Environment Setup

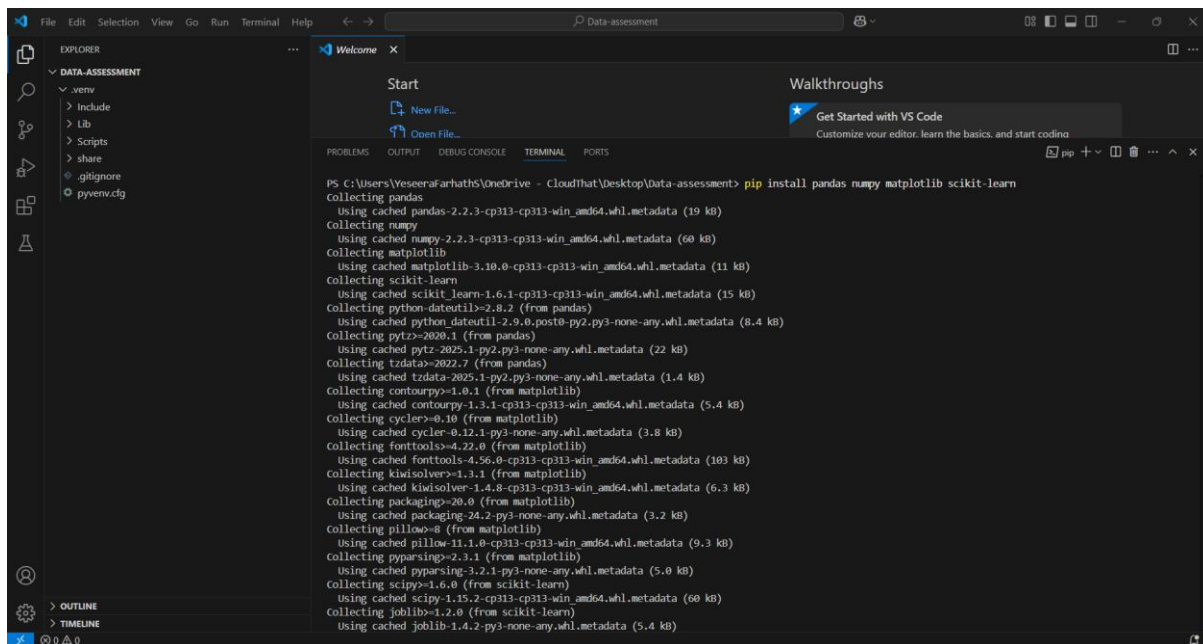
- Use VSCode Mandatory
- Create project structure (data, notebooks, scripts, output folders)
- Set up Jupyter Notebooks

**Solution:****I) Team Setup:****1. Set up virtual environments for Python dependencies**

- Created a new folder
- Created a python virtual environment.

**2. Create a new GitHub repository for the project**

### 3. Install required Python libraries (NumPy, Pandas, Matplotlib, scikit-learn)

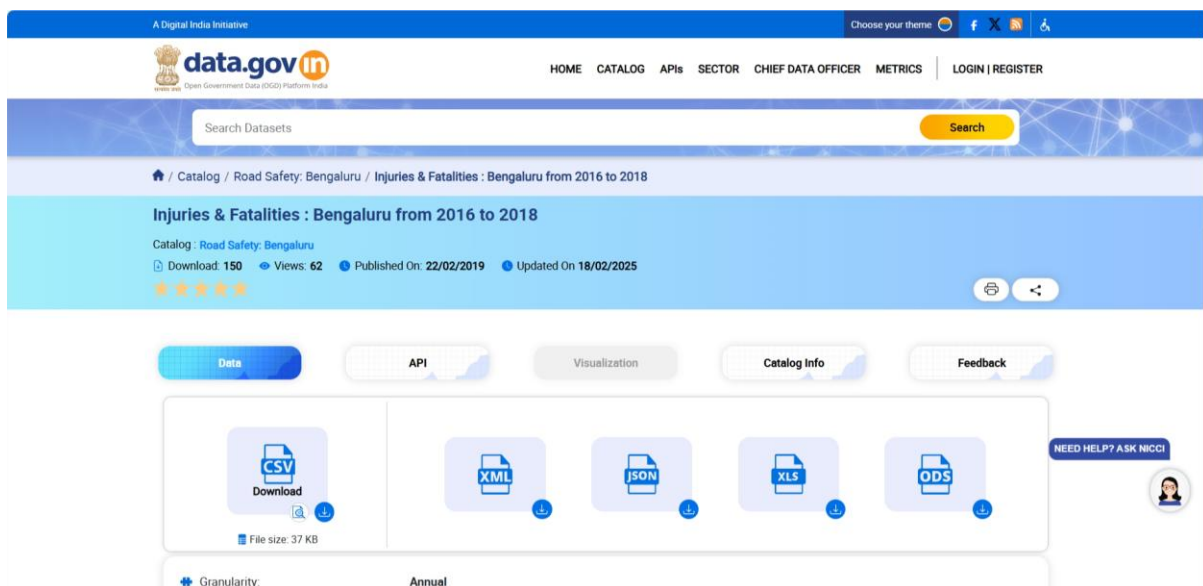


```

PS C:\Users\YeseeraFarhath\OneDrive - Cloudhat\Desktop\Data-assessment> pip install pandas numpy matplotlib scikit-learn
Collecting pandas
  Using cached pandas-2.2.3-cp313-cp313-win_amd64.whl.metadata (19 kB)
Collecting numpy
  Using cached numpy-2.2.3-cp313-cp313-win_amd64.whl.metadata (60 kB)
Collecting matplotlib
  Using cached matplotlib-3.10.0-cp313-cp313-win_amd64.whl.metadata (11 kB)
Collecting scikit-learn
  Using cached scikit-learn-1.6.1-cp313-cp313-win_amd64.whl.metadata (15 kB)
Collecting python-dateutil<=2.8.2 (from pandas)
  Using cached python-dateutil-2.9.0.post0-py2.py3-none-any.whl.metadata (8.4 kB)
Collecting pytz<=2020.1 (from pandas)
  Using cached pytz-2025.1-py2.py3-none-any.whl.metadata (1.4 kB)
Collecting tzdata<=2025.1-py2.py3-none-any.whl.metadata (1.4 kB)
Collecting contourpy>=1.0.1 (from matplotlib)
  Using cached contourpy-1.3.1-cp313-cp313-win_amd64.whl.metadata (5.4 kB)
Collecting cycler>=0.10 (from matplotlib)
  Using cached cycler-0.12.1-py3-none-any.whl.metadata (3.8 kB)
Collecting fonttools>=4.22.0 (from matplotlib)
  Using cached fonttools-4.56.0-cp313-cp313-win_amd64.whl.metadata (103 kB)
Collecting kiwisolver>=1.3.1 (from matplotlib)
  Using cached kiwisolver-1.4.8-cp313-cp313-win_amd64.whl.metadata (6.3 kB)
Collecting packaging>=20.0 (from matplotlib)
  Using cached packaging-24.2-py3-none-any.whl.metadata (3.2 kB)
Collecting pillow>=8 (from matplotlib)
  Using cached pillow-11.1.0-cp313-cp313-win_amd64.whl.metadata (9.3 kB)
Collecting pyparsing>=2.3.1 (from matplotlib)
  Using cached pyparsing-3.2.1-py3-none-any.whl.metadata (5.0 kB)
Collecting scipy>=1.6.0 (from scikit-learn)
  Using cached scipy-1.15.2-cp313-cp313-win_amd64.whl.metadata (60 kB)
Collecting joblib>=1.2.0 (from scikit-learn)
  Using cached joblib-1.4.2-py3-none-any.whl.metadata (5.4 kB)

```

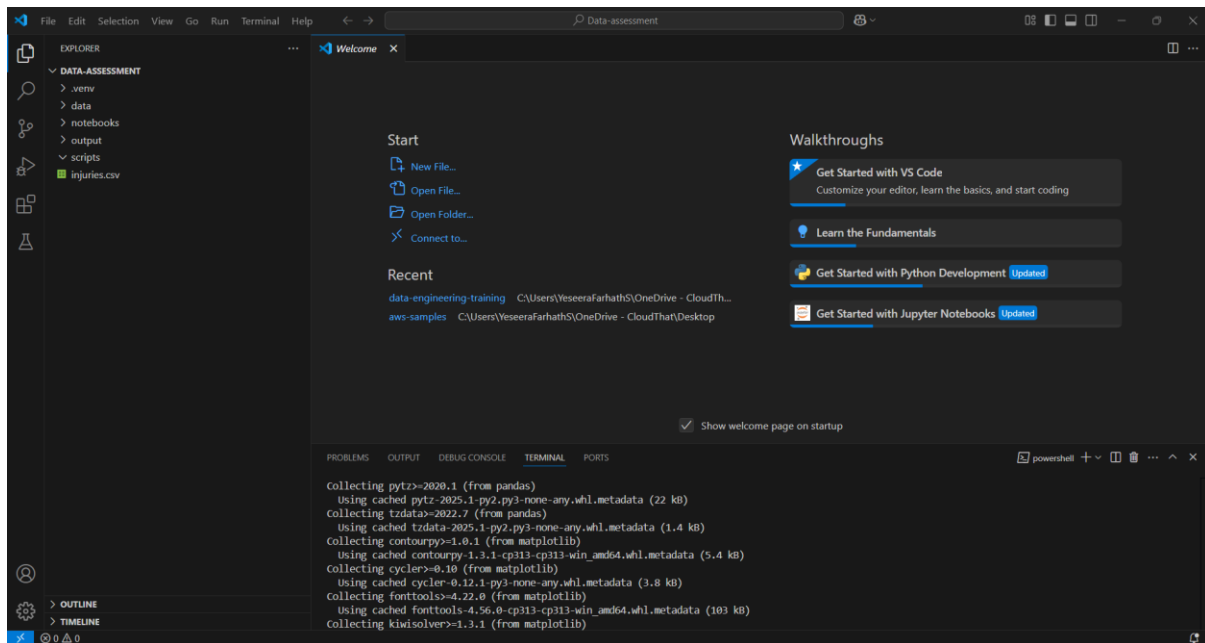
### II) Dataset Selection: Choose any from one dataset from data.gov.in



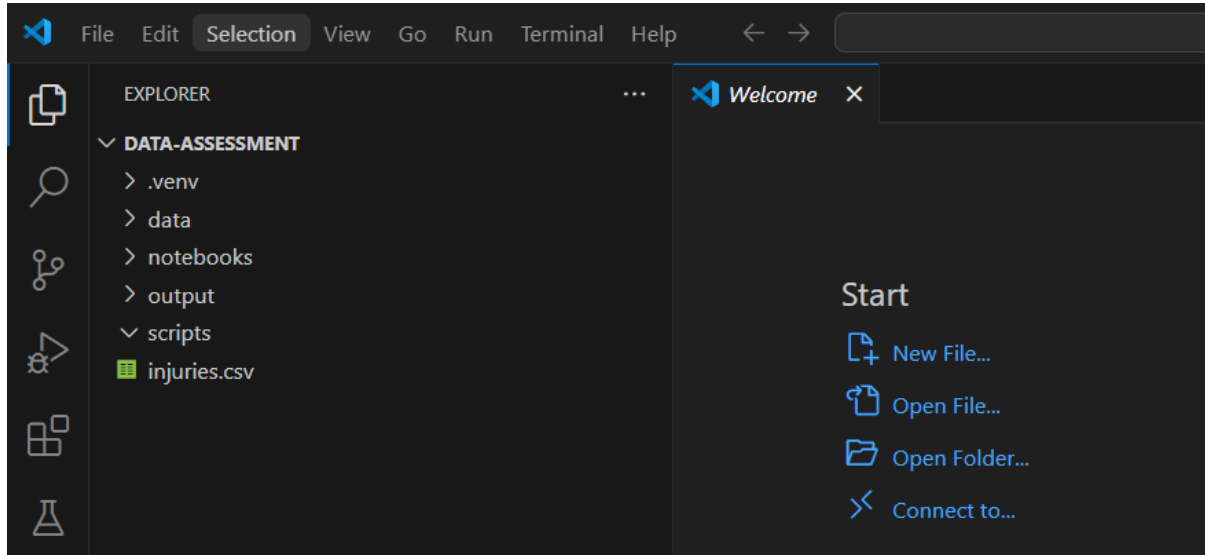
- Selected Injuries and fatalities dataset

### III) Environment Setup

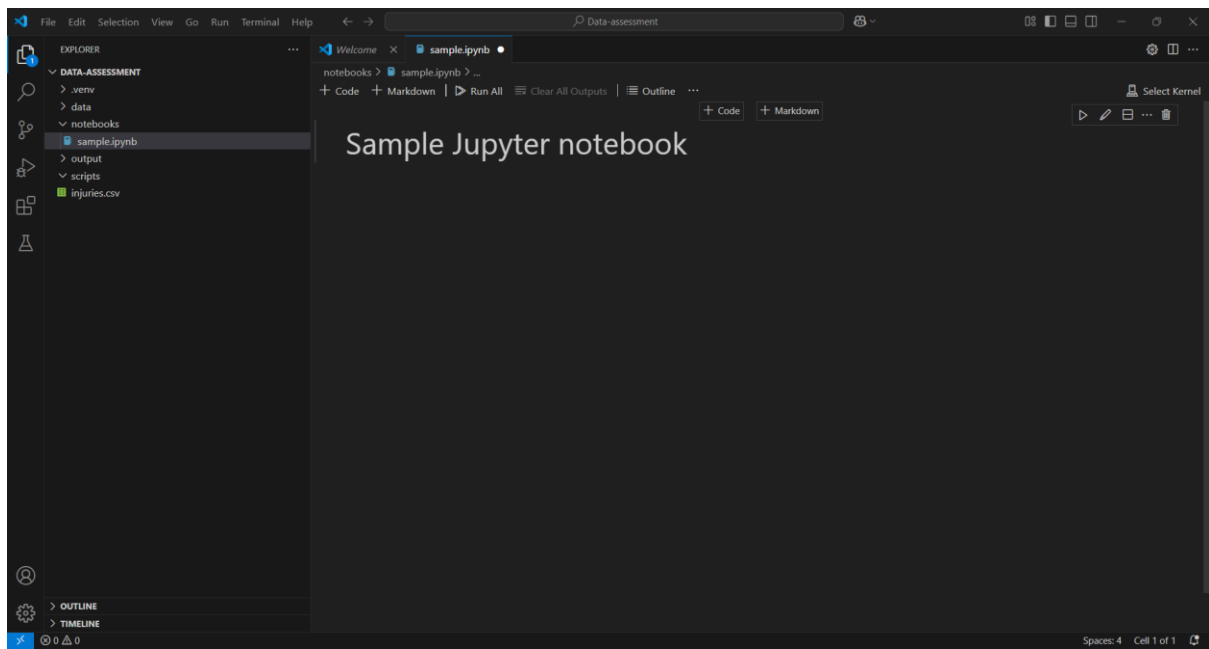
#### 1. Use VSCode Mandatory



#### 2. Create project structure (data, notebooks, scripts, output folders)



### 3. Set up Jupyter Notebooks



## Day 2: Data Processing and Analysis

- **1. Data Processing**
  - o Use Pandas to:
    - § Clean the dataset (handle missing values, format dates)
    - § Create derived features relevant to the chosen dataset
    - § Aggregate data for meaningful analysis
- **2. Exploratory Analysis**
  - o Use Jupyter Notebooks to:
    - § Generate descriptive statistics
    - § Create visualizations using Matplotlib
    - § Document key findings and insights
    - § Explore relationships between variables

Solution:

1. Import required packages:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

[2] ✓ 7.2s

## 2. Import the dataset (injuries.csv)

```

df = pd.read_csv('injuries.csv')
df.info()

```

[29] ✓ 0.0s

```

... <class 'pandas.core.frame.DataFrame'>
RangeIndex: 963 entries, 0 to 962
Data columns (total 32 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   City Name                                                             1 non-null     object
1   2018 - Total Injuries - Pedestrian                                   1 non-null     float64
2   2018 - Total Injuries - Bicycles                                     1 non-null     float64
3   2018 - Total Injuries - Two-wheelers                                 1 non-null     float64
4   2018 - Total Injuries - Other modes of road transport (auto, bus, lorry) 1 non-null     float64
5   2018 - Total Injuries                                                1 non-null     float64
6   2018 - Total Fatalities - Pedestrian                                 1 non-null     float64
7   2018 - Total Fatalities - Bicycles                                   1 non-null     float64
8   2018 - Total Fatalities - Two-wheelers                               1 non-null     float64
9   2018 - Total Fatalities - Other modes of road transport (auto, bus, lorry) 1 non-null     float64
10  2018 - Total Fatalities                                               1 non-null     float64
11  2017 - Total Injuries - Pedestrian                                    1 non-null     float64
12  2017 - Total Injuries - Bicycles                                     1 non-null     float64
13  2017 - Total Injuries - Two-wheelers                                 1 non-null     float64
14  2017 - Total Injuries - Other modes of road transport (auto, bus, lorry) 1 non-null     float64
15  2017 - Total Injuries                                                1 non-null     float64
16  2017 - Total Fatalities - Pedestrian                                 1 non-null     float64
17  2017 - Total Fatalities - Bicycles                                   1 non-null     float64
18  2017 - Total Fatalities - Two-wheelers                               1 non-null     float64
19  2017 - Total Fatalities - Other modes of road transport (auto, bus, lorry) 1 non-null     float64
...
30  2016 - Total Fatalities.1                                             1 non-null     float64
31  Unnamed: 31                                                           0 non-null     float64
dtypes: float64(31), object(1)
memory usage: 240.9+ KB

```

- Using Toyota.csv dataset, as injuries.csv is almost null.

```

df = pd.read_csv('Toyota.csv')
df

```

[3] ✓ 0.0s

```

...

```

|      | Unnamed: 0 | Price | Age  | KM    | FuelType | HP  | MetColor | Automatic | CC   | Doors | Weight |
|------|------------|-------|------|-------|----------|-----|----------|-----------|------|-------|--------|
| 0    | 0          | 13500 | 23.0 | 46986 | Diesel   | 90  | 1.0      | 0         | 2000 | three | 1165   |
| 1    | 1          | 13750 | 23.0 | 72937 | Diesel   | 90  | 1.0      | 0         | 2000 | 3     | 1165   |
| 2    | 2          | 13950 | 24.0 | 41711 | Diesel   | 90  | NaN      | 0         | 2000 | 3     | 1165   |
| 3    | 3          | 14950 | 26.0 | 48000 | Diesel   | 90  | 0.0      | 0         | 2000 | 3     | 1165   |
| 4    | 4          | 13750 | 30.0 | 38500 | Diesel   | 90  | 0.0      | 0         | 2000 | 3     | 1170   |
| ...  | ...        | ...   | ...  | ...   | ...      | ... | ...      | ...       | ...  | ...   | ...    |
| 1431 | 1431       | 7500  | NaN  | 20544 | Petrol   | 86  | 1.0      | 0         | 1300 | 3     | 1025   |
| 1432 | 1432       | 10845 | 72.0 | ??    | Petrol   | 86  | 0.0      | 0         | 1300 | 3     | 1015   |
| 1433 | 1433       | 8500  | NaN  | 17016 | Petrol   | 86  | 0.0      | 0         | 1300 | 3     | 1015   |
| 1434 | 1434       | 7250  | 70.0 | ??    | NaN      | 86  | 1.0      | 0         | 1300 | 3     | 1015   |
| 1435 | 1435       | 6950  | 76.0 | 1     | Petrol   | 110 | 0.0      | 0         | 1600 | 5     | 1114   |

1436 rows × 11 columns

## 3. Get information about the dataset

```
df.info()
```

[4] ✓ 0.0s

```
... <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1436 entries, 0 to 1435
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   1436 non-null   int64
1   Price        1436 non-null   int64
2   Age          1336 non-null   float64
3   KM           1436 non-null   object
4   FuelType     1336 non-null   object
5   HP           1436 non-null   object
6   MetColor     1286 non-null   float64
7   Automatic    1436 non-null   int64
8   CC           1436 non-null   int64
9   Doors        1436 non-null   object
10  Weight       1436 non-null   int64
dtypes: float64(2), int64(5), object(4)
memory usage: 123.5+ KB
```

## 4. Dropping the unwanted column named 'unnamed'

```
df.drop(columns=['Unnamed: 0'], inplace=True, errors='ignore')
```

[5] ✓ 0.0s

```
df.columns
```

[6] ✓ 0.0s

```
... Index(['Price', 'Age', 'KM', 'FuelType', 'HP', 'MetColor', 'Automatic', 'CC',
        'Doors', 'Weight'],
        dtype='object')
```



## 5. Check for null values

```

▶ df.isnull().sum()
[7] ✓ 0.0s

... Price      0
    Age      100
    KM        0
    FuelType  100
    HP         0
    MetColor  150
    Automatic  0
    CC         0
    Doors      0
    Weight     0
    dtype: int64

```

## 6. Clean the data (missing values, wrong data types, and question mark)

```

▶ # Replace ?? with NaN
df.replace("?", np.nan, inplace=True)

# Convert data types
df['KM'] = df['KM'].str.replace(',', '').astype(float)
df['HP'] = pd.to_numeric(df['HP'], errors='coerce')
df['Doors'] = df['Doors'].replace({'three': 3, 'four': 4, 'five': 5}).astype(float)

# missing values
df.fillna({
    'Age': df['Age'].median(),
    'KM': df['KM'].median(),
    'HP': df['HP'].median(),
    'FuelType': df['FuelType'].mode()[0],
    'MetColor': df['MetColor'].mode()[0],
    'Doors': df['Doors'].mode()[0]
}, inplace=True)

[8] ✓ 0.0s

df.isnull().sum()
[9] ✓ 0.0s

... Price      0
    Age      0
    KM        0
    FuelType   0
    HP         0
    MetColor   0
    Automatic   0
    CC         0
    Doors      0
    Weight     0
    dtype: int64

```

## 7. Plot graphs of clean data

```
# Plot distributions of numeric features
fig, axes = plt.subplots(3, 2, figsize=(12, 12))
fig.suptitle("Distributions of Key Numeric Features", fontsize=14)

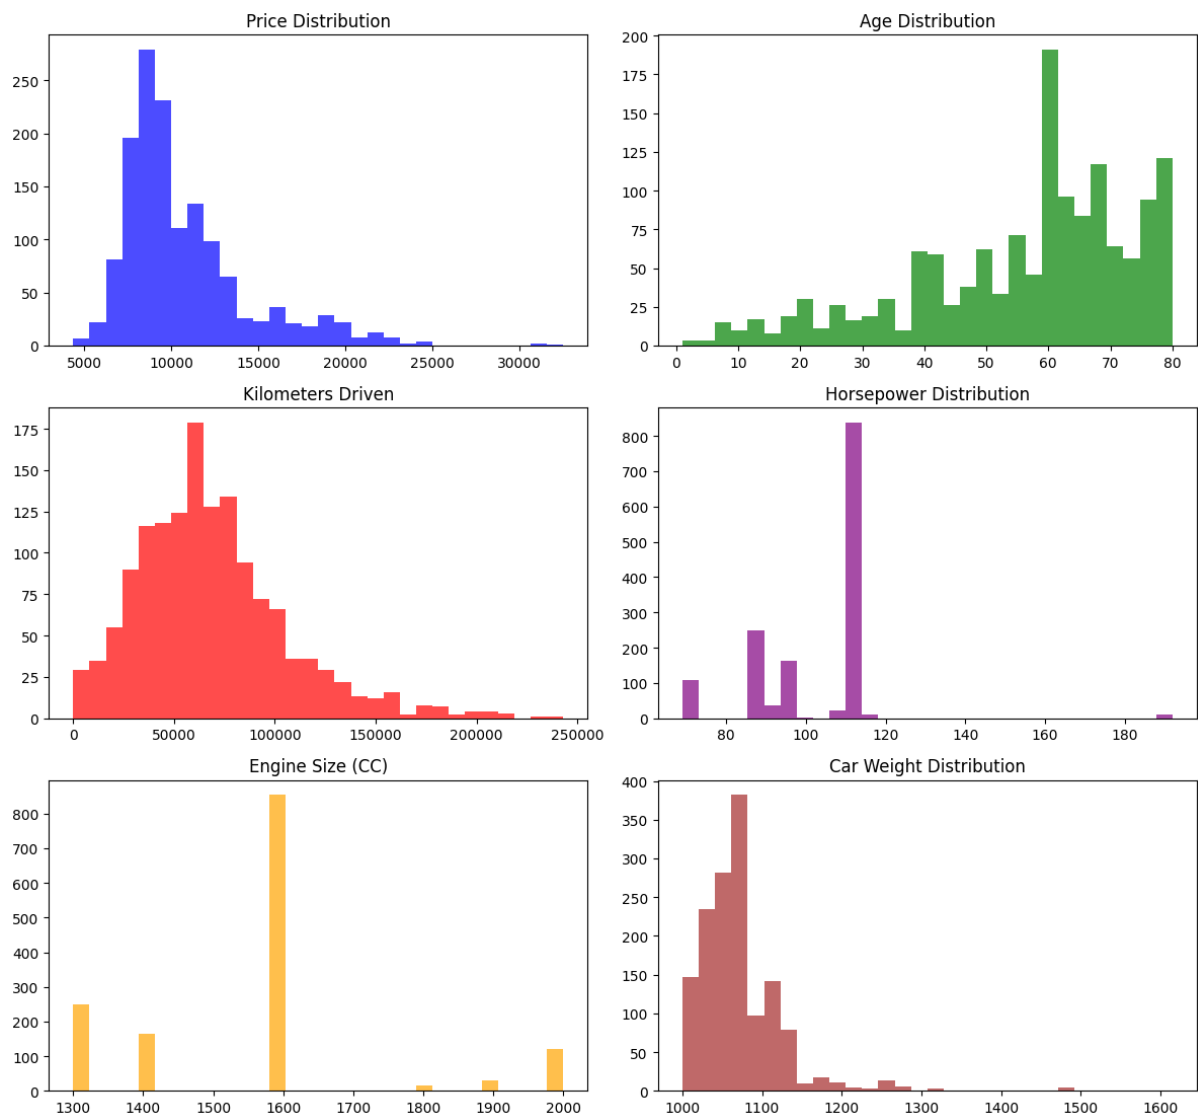
axes[0, 0].hist(df['Price'], bins=30, color='blue', alpha=0.7)
axes[0, 0].set_title("Price Distribution")
axes[0, 1].hist(df['Age'], bins=30, color='green', alpha=0.7)
axes[0, 1].set_title("Age Distribution")
axes[1, 0].hist(df['KM'], bins=30, color='red', alpha=0.7)
axes[1, 0].set_title("Kilometers Driven")
axes[1, 1].hist(df['HP'], bins=30, color='purple', alpha=0.7)
axes[1, 1].set_title("Horsepower Distribution")
axes[2, 0].hist(df['CC'], bins=30, color='orange', alpha=0.7)
axes[2, 0].set_title("Engine Size (CC)")
axes[2, 1].hist(df['Weight'], bins=30, color='brown', alpha=0.7)
axes[2, 1].set_title("Car Weight Distribution")

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()
```

[11] ✓ 1.4s


## 8. Plots of numeric features

Distributions of Key Numeric Features



## 9. Fuel type distribution

```
# Fuel Type Distribution
plt.figure(figsize=(8, 5))
df['FuelType'].value_counts().plot(kind='bar', color=['teal', 'yellow', 'red'])
plt.title("Fuel Type Distribution")
plt.xlabel("Fuel Type")
plt.ylabel("Count")
plt.show()
```



## 10. Plots for relationships between features


```
# Scatter plots for relationships between key variables
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

axes[0].scatter(df["Age"], df["Price"], alpha=0.3, color='blue')
axes[0].set_title("Price vs. Age")
axes[0].set_xlabel("Age (months)")
axes[0].set_ylabel("Price")

axes[1].scatter(df["KM"], df["Price"], alpha=0.3, color='red')
axes[1].set_title("Price vs. Kilometers Driven")
axes[1].set_xlabel("KM Driven")
axes[1].set_ylabel("Price")

plt.tight_layout()
plt.show()
```

[17] ✓ 0.3s



## 11. Insights

- **Price vs. Age:** Older cars have lower prices
- **Price vs. KM Driven:** Cars with higher kilometers driven have lower prices.
- **Fuel Type Popularity:** Petrol cars are the most common, followed by diesel, with very few LPG cars in the dataset.
- **Horsepower & Price:** Cars with higher horsepower have higher prices.

Day -3

### Day 3: Modeling and Presentation

- **1. Predictive Modeling**
  - o Build an appropriate predictive model based on the dataset chosen
  - o Evaluate model performance using appropriate metrics
  - o Document model selection process and rationale
- **2. Documentation and Presentation**
  - o Document the entire process
  - o Prepare presentation slides
  - o Create a final report with actionable insights

Solution:

1. Import sklearn and define features and target

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

features = ['Age', 'KM', 'HP', 'CC', 'Weight']
target = 'Price'
```

[20] ✓ 18.0s

2. Split the dataset for training and testing

```
# Split data
X = df[features]
y = df[target]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

[22] ✓ 0.0s

3. Choose and apply model – Linear regression model is selected

```
▶ # linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

[23] ✓ 0.0s
```

... LinearRegression ⓘ ?  
LinearRegression()

4. Make predictions

```
# Make predictions
y_pred = model.predict(X_test)

[24] ✓ 0.0s
```

5. Evaluate the model using metrics

```
▶ # Metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error: {mae}")
print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")

[25] ✓ 0.0s
```

... Mean Absolute Error: 1007.7523623460324  
Mean Squared Error: 2467859.499716883  
R-squared: 0.8150417031532511

## 6. Plot of actual vs predicted price

```
# Plot actual vs. predicted prices
plt.figure(figsize=(8, 5))
plt.scatter(y_test, y_pred, alpha=0.5, color='blue')
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual vs. Predicted Prices")
plt.show()
```

